

An Ethereum Light Client on Axelar

Shresth Agrawal^{1,2}, Dionysis Zindros^{1,3},
Dimitris Karakostas^{1,4}, and Apostolos Tzinias^{1,5}

¹ Common Prefix

² Technische Universität München

³ Stanford University

⁴ University of Edinburgh

⁵ National Technical University of Athens

May 1, 2023

Last update: December 14, 2025

Abstract. Axelar network is a decentralized interoperability layer that connects a large set of blockchains. Axelar is based on the Cosmos SDK and defines a set of Tendermint proof-of-stake validators. These validators play a significant role in relaying events from supported chains to the Axelar network and are expected to run full nodes of the supported EVM chains to gather event information. However, this operation is typically expensive, such that the entry and maintenance costs for validators increase, along with possibly leading some validators to use centralized RPC providers. The newly launched Axelar Interchain Amplifier protocol allows us to easily swap out this functionality with a more robust solution. This work aims to reduce operational costs and increase Axelar’s level of decentralization. Specifically, we advocate for an on-chain light client to run as CosmWasm contracts on Axelar, integrated into the new Amplifier model, that trustlessly consumes Ethereum chain data. Such a light client checks the veracity of cross-chain claims to ensure the relevant events have been included in the canonical source chain, safeguarding against untrustworthy EVM RPC providers. Focusing on Ethereum, we discuss and contrast several possible flavors of light or superlight clients based on *Simple Payment Verification* (SPV), *zero knowledge*, or *refereed games*. Tailoring the solution to the particularities of Axelar, we recommend SPV as the most suitable choice. We present an architecture blueprint and estimate the workforce allocation. Lastly, we put forth a theoretical model to support the need for such design changes. Theoretically, this work will enable the system to endure temporary dishonest advantages among Axelar validators and restore safety following the re-establishment of an honest supermajority.

1 Introduction

Blockchain technology has revolutionized the way we create trustless and decentralized applications, offering immense composability and the ability to build complex primitives. However, the blockchain landscape is far from homogeneous, with multiple ecosystems that are based on their own consensus protocols and offer varying security and operation features.

Axelar⁶ enables interoperability between diverse blockchain ecosystems. Axelar is based on the Cosmos SDK⁷ and aims to interlink various blockchain ecosystems, such as Cosmos, Ethereum, Bitcoin, Polygon, and others. Notably, it has been identified as one of only two possible solutions that are sufficiently secure across the stack to be used by the leading decentralized exchange, Uniswap [13].

Ethereum, in particular, is one of the most widely used blockchains, hosting a multitude of decentralized applications spanning various sectors, including finance, gaming, and governance [3,14]. Therefore, establishing a seamless and secure connection to Ethereum is of paramount importance for Axelar. This work focuses on building a trustless connection from Ethereum to Axelar.

⁶<https://axelar.network>

⁷<https://cosmos.network>

1.1 Existing Construction

Axelar implements a Tendermint-based delegated proof-of-stake consensus mechanism [2, 10]. In particular, validators lock stake in the Axelar chain and receive stake delegations from Axelar users. The top 75 validators, based on the aggregate (self and delegated) stake, are chosen to participate in Axelar’s consensus mechanism.

Axelar acts as the hub connecting different source chains. Every connection with Axelar consists of two essential components. The first component verifies source chain (e.g., Ethereum) data into Axelar. The second component generates threshold signatures on Axelar, which can be verified on the source chain. In this report, we focus exclusively on the former, that is verification of source chain data into Axelar.

An on-chain voting mechanism within Axelar is used to verify transactions that occur on the source chain. This process is carried out by a set of *attestors*, who are required to run a full node of the source chain and have access to the full node’s RPC (Remote Procedure Call) interface. This enables them to verify the finalized transactions on the source chain, before voting to bridge them on Axelar.

To bridge data from the source chain to Axelar, a user interacts with an Axelar smart contract on the source chain. Subsequently, a voting poll is initiated on Axelar, which is viewed by attestors. For each poll, an attester decides whether to vote for or against. To make an informed decision, the attester queries the source chain’s full node RPC and checks if the transaction under question has been finalized on the source chain. If a poll receives sufficient attestations, it is accepted, otherwise it is rejected.

This voting process forms the basis of verifying the source chain data into Axelar. Nonetheless, for the connector construction to function securely, both the source chain and Axelar are assumed safe and live. In addition, the voting power distribution among the attestors is assumed to have an honest majority.

1.2 Amplifier Construction

The new Interchain Amplifier construction [11] generalizes Axelar’s bridging protocol and brings it to the execution layer, introducing a programmable verification interface through which new chains can integrate with Axelar.

The verification logic is now implemented through CosmWasm smart contracts, allowing each source chain connection to define its own verification method. This enables alternative, trust-minimized verification paths, such as light clients, alongside or in place of the default voting attestation flow. Such solutions can eliminate the operational overhead of attestors monitoring external chains, participating in attestation polls, or operating additional infrastructure. Consequently, this allows to reduce the overall cost of securing cross-chain connections, while strengthening their security at the same time.

1.3 Problem Statement

The current construction of Axelar relies on attestors running the full node of the source chain to verify transactions and vote in an informed manner. However, running a full node, or many, imposes significant costs on attestors.⁸ In order to maintain a diverse and robust ecosystem, comprising a multitude of participants, it is important to develop cost-effective

⁸Indicatively, running a full Ethereum node requires 4+ CPU cores, 16+ GB RAM, at least 1 TB SSD, and 25 Mbps of stable connection (source: <https://www.quicknode.com/guides/infrastructure/node-setup/ethereum-full-node-vs-archive-node>).

and efficient solutions.⁹ Such solutions could reduce the entry and maintenance costs for new validators, also possibly leading to more value being gained for users.

To address this challenge, we propose a construction that enables verification of the source chain’s consensus within the Axelar execution layer. This is accomplished through light and super-light client constructions of the source chain, implemented as CosmWasm smart contracts on Axelar, and integrated into the Interchain Amplifier protocol.

In this report, we explore several constructions for light and super-light clients tailored specifically for Ethereum. We then propose a construction that best aligns with Axelar’s vision and requirements, aiming to drive decentralization, enhance security, and improve scalability. Our construction makes use of Ethereum’s sync committee and guarantees bridge safety.




Fig. 1: We recommend implementing an on chain light client using sync committee. The light client will be implemented as CosmWasm smart contracts deployed on the Axelar network. The light client will be responsible for verifying the handover signatures, block signatures and Merkle inclusion proofs.

1.4 Preliminaries

Bridge

A *bridge* is an interoperability protocol between two ledger protocols. The purpose of the bridge is to relay events or information that take place on the source side to the destina-

⁹In practice, systems that impose significant costs for running a full node often demonstrate centralization tendencies, where participants tend to rely on third-party service providers to run and maintain the full nodes on their behalf. Due to economies of scale, a large portion of these networks tends to cluster around a small number of such providers; e.g., at times, 50% of Ethereum’s transactions ran through one such provider, Infura [8]. Therefore, proactively addressing such hazards can lead to a healthier, more diverse and robust ecosystem.

tion side [9, 16]. In particular, the parties that maintain the bridge transmit *cross-chain* transactions, such that a transaction on the source side is represented by an “image” on the destination side. There are two core properties that a secure bridge should guarantee: safety and liveness.¹⁰

Bridge Safety. Bridge safety mandates that a transaction appears on the destination side *only if* it has first appeared on the source side, albeit with some delay. Intuitively, safety ensures that *bad things don’t happen*, that is it’s impossible to find a transaction paying out on the destination without its “pre-image” having appeared on the source side. In essence, if safety is guaranteed, an adversary cannot create money on the destination side without having paid on the source side.

Bridge Liveness. Bridge liveness ensures that a transaction which appears on the source side will eventually make it to the destination side. Intuitively, this guarantees that *good things happen*, that is whenever an honest party attempts to cross the bridge, it will successfully do so.

Light Clients

A *light client* is a client that wishes to synchronize with the rest of the network, but has limited resources available in terms of communication, computation, and storage. One example of a limited resource computer is the on-chain smart contract infrastructure of Axelar, where computation and storage are expensive. The light client begins its lifecycle holding the genesis block and synchronizes to the current tip of the canonical chain once in a while. Our job when building a light client is, given a block that the light client has already downloaded sometime in the past, to allow the client to synchronize with the most recent chain tip.

Ethereum Sync Committee

One useful ingredient of the Ethereum ecosystem that enables the construction of efficient light clients is the so-called *sync committee* [6]. This feature was introduced in the system’s “Altair” hard fork and is specifically tailored for use of light clients, as we will discuss in the alternatives of Section 2.

The sync committee consists of 512 of Ethereum validators. They are randomly selected, from the set of all validators, every sync committee period (approx. 1 day). Each honest validator in the sync committee is continually online and signs the header of each block during the sync committee period.

The sync committee of period X is decided one periods in advance, that is the beginning of period $X - 1$. As headers of the current period X containing the root of the next sync committee $X + 1$ are signed by the current sync committee X , the signatures of the current sync committee act as a handover mechanism from the current committee to the next. A prover can provide signatures from $2/3$ of the current committee on some header of current period as a proof of validity of the next committee. Note that there can be multiple headers in the current period which can have more than $2/3$ sync committee signatures on them. The prover can choose any of these headers with their respective signatures as a proof.

We demonstrate the usage of the sync committee with the following example. Assume that a client C holds the block header of some slot N , that is part of the sync committee period X . When C wants to authenticate the header of a block at slot N' , which is part of

¹⁰Appendix 5.1 offers a formal definition of these properties.

the period $X + 1$, it proceeds as follows. First, C validates the sync committee of period $X + 1$. To do so, it verifies that the Merkle root of the committee was published in a block header of period X ; if so, C updates its sync committee for $X + 1$ as the one defined in the Merkle tree. At this point, C can validate every block header for each slot of period $X + 1$, by obtaining the corresponding signatures of the sync committee of that period.

Using this iterative process, the light client can validate all future blocks, starting from an initial trusted point. The cost of validating a block depends on the size of the committee, the aggregate signature, and the Merkle path; in Ethereum this is estimated to approx. 25KB [6].

There are two main points of discussion around the sync committee. First, the committee should be honest. In essence, the sampling process, via which the sync committee's members are chosen from the set of all Ethereum validators, should be secure, s.t. the probability that $\frac{2}{3}$ of the committee members are adversarial is negligible. Observe that, if the committee is malicious, then it can convince a light client of a fraudulent Ethereum state. Second, slashing is currently not implemented in the sync committee. Therefore, a committee member can double-sign, i.e., sign conflicting blocks, without direct counterincentives. Whether committee members are incentivized to behave in such manner is outside the scope of this document, although it is an interesting research question.

2 On-Chain Light Client

We now present an array of solutions for the on-chain light client. For each solution we give a high-level overview, highlighting its mechanics as well as possible shortcomings. Each alternative offers different features and relies on different assumptions.

We note that all solutions, except the first, rely on Ethereum's sync committee.

SPV Light Client (sampling validators)

The first solution makes use of Ethereum's validators. In particular, we assume that there exists a mechanism for sampling some of the signatures that were produced by the validators on each block. To verify an Ethereum block, the light client checks: i) that the sampling process was done correctly (e.g., following a verifiable pseudorandom process); ii) that the provided signatures correspond to active Ethereum validators; iii) that enough signatures have been provided to guarantee the block's correctness.

An advantage of this solution is the non-reliance on the sync committee. In particular, the only assumption that is needed is that Ethereum is safe and live (which is already assumed) and that the probability of sampling adversarial signatures is negligible.

However, it also presents various shortcomings. First, Ethereum's consensus relies on Casper FFG [4] and LMD-GHOST. This is particularly complicated and it is unclear how a secure signature sampling process can be implemented. Second, it is unclear how the light client can retrieve the active Ethereum validators at any point in time. In particular, assuming that the light client starts from a trusted block header, it is unclear how to obtain each following period's validators, in a succinct and efficient manner.

SPV Light Client (Sync committee)

This solution makes direct use of Ethereum's sync committee. In particular, each sync committee is recorded on Axelar. When an attestor wants to bridge a transaction from Ethereum to Axelar, it accompanies it with a proof of inclusion in a block that has been validated by the corresponding sync committee. Following, when an Axelar full nodes

wants to validate an Ethereum transaction that has been bridged to Axelar, it first identifies (on Axelar’s chain) the sync committee and that corresponds to the transaction and then validates the proof w.r.t. it.

This scheme assumes that, at the bridge’s onset, the sync committee is recorded on Axelar correctly. In other words, the first sync committee that is recorded on Axelar should be correct. Following, each update to the sync committee is accompanied by a proof of a handover, that is an Ethereum block that has been signed by the previous committee. Note that, even if Axelar’s security is compromised, the adversary cannot produce a proof of a handover to an invalid sync committee (since this depends on the sync committee’s security). Therefore, the honest Axelar nodes will always hold a valid list of sync committees (albeit this could possibly be outdated, if liveness is violated).

The major advantage of this solution is that it is straightforward to implement. In particular, the sync committee has already been implemented in Ethereum and there exists a large body of community projects and tools that already implement light client based on sync committee, such as Kevlar¹¹ and Helios¹².

There are some disadvantages with this solution though. First, it relies on the security of Ethereum’s sync committee. Second, the storage requirements increase linearly over time, since each sync committee (which is updated approx. every day) is recorded on Axelar’s chain. This can result in significant storage overhead, since the mechanism requires approx. 25KB per committee [6].

Refereed Light Client

This mechanism makes use of Ethereum’s sync committee in a manner similar as the previous solution, but aiming to reduce the storage requirements.

In brief, the idea is to (optimistically) assume that an attester provides honest data about the updated sync committees and enable a dispute resolution mechanism in case the attester is malicious [1]. Specifically, we assume that the bridge is updated every x (sync committee) periods. In the previous solution, each committee is recorded on Axelar’s chain. Here, there exists an Axelar attester who collects the x committees, each corresponding to the sync committees for the corresponding periods, and publishes a commitment to them on Axelar altogether. If the attester is honest, then the commitment should be correct. Therefore, the other Axelar nodes can “fast-forward” these x committees.

Nonetheless, the attester might be malicious and submit an incorrect commitment. To cover this possibility, there exists a contest period, during which another attester can challenge the first attester’s submission. The challenger provides an alternative commitment, at which point the dispute needs to be resolved. This is done by opening both commitments and identifying the point of divergence between the two committee lists. At this point, both attestors are required to submit proof of correct transition, that is to reveal the (sync committee) keys that correspond to the last agreed committee along with (this committee’s) signatures on the keys that correspond to the first committee of disagreement. Since the sync committee is presumed honest, one of the two lists will be revealed as fraudulent, since it should be impossible to provide the necessary signatures that validate the transition under question.

We propose two manners in which the commitment can be implemented.

Linear In this case, the commitment is a list of hashes, with each hash corresponding to a sync committee. The main benefit of this solution is that disputes can be resolved easily,

¹¹<https://github.com/lightclients/kevlar>

¹²<https://github.com/a16z/helios>

since the point of disagreement can be identified directly by comparing the two lists element by element. Once the point of disagreement is identified the attestors need to open the last agreed committee, the committee of disagreement and the handover signatures from the last agreed committee to the committee of disagreement. However, the asymptotic complexity of this solution is again linear on the number of committees, the proof size is significantly smaller than the previous solution as only committee hashes have to be submitted instead of complete sync committee public keys.

Bisection Here, the commitment is a Merkle tree, the leaves of which correspond to the hash of each sync committee. Initially the attester publishes the root of the tree. In case of a dispute, the challenger provides an alternative root and initiates an interactive dispute resolution protocol. In each round the attestors open the children of the disputed node. Then the on chain protocol checks if the children are correct (i.e. the hash of the children matches the hash of the node) and compares left child and then the right child. If the left child is already different for the two attestors the next disputed node is set to the left child, otherwise to the right child. This process is recursively repeated until we reach the leaf of the tree which is also the first point of disagreement. At this point the attestors can follow the same protocol as the linear case to resolve the dispute.

The advantage of this solution is that the storage complexity is constant in the optimistic case (i.e. the attestors are honest and no dispute resolution is required) where as it falls back to poly logarithmic complexity in the case of dispute resolution. [1, 12].

Finally, we note that the advantage of the refereed light client solution only manifests if the bridge is synchronized infrequently and dispute resolutions are rare. If, for example, the bridge is synchronized every day (that is, $x = 1$), then the benefit is marginal even in the optimistic case. In addition, implementing a dispute resolution mechanism securely adds complexity. Finally, the length of the contest period affects the bridge's latency, since it should be large enough s.t. an honest party can challenge an invalid commitment, but small enough to avoid increasing latency significantly.

SNARK Light Client

The final solution is based on zero-knowledge, namely SNARK proofs (Succinct Non-interactive ARguments of Knowledge). The goal of this solution is to reduce the proof of validity of each bridged Ethereum transaction.

The setup of this solution is similar to the sync committee-based solution above. Specifically, each committee hash is published on Axelar, along with a proof of handover for the committee transitions.

In particular, when an attester wants to bridge an Ethereum transaction to Axelar, it produces a zero-knowledge proof. This proof replaces the aggregate public key and signatures, which were used before for validation. Instead, it proves that there exists a quorum of at least $\frac{2}{3}$ sync committee members that have signed the transaction under question.

The main advantage of this solution is that SNARKs can be very efficient in terms of proof sizes. For example, a Groth16 [7] proof consists of only 2 group elements, that is approx. 209 bytes. In fact we can generate a constant size proof for increasing length of sync periods. But this benefit can not be realised in the current bridge construction because the bridge is continuously synchronised and therefore the proof data submitted to the chain will grow linearly with the number of sync periods.

One main disadvantage of this solution is implementation complexity and cost. Securely implementing ZK-SNARKs can be particularly challenging, even more so since they

have been developed fairly recently. Additionally, the cost of generating a SNARK can be particularly high. For example, according to [15], proving consensus (of 128 validators) on *one* Cosmos block takes 18 seconds on 32 instances of Amazon AWS `c5.24xlarge`. (We are unable to independently reproduce this finding because the authors of zkBridge have not disclosed their code.) At Cosmos’ block rate of 1 block per second, that would require 18×32 continuously operating `c5.24xlarge` instances, costing annually \$12,967,488 on Amazon AWS (annual pricing, `us-east-1` region, June 2023), or \$1,749,600 on Hettner’s equivalents. Scaling this proportionally for Ethereum yields an estimate of \$583,333 annually.

Comparison

Table 1 summarizes the comparison of the alternatives presented above. Note that, in all cases, Ethereum is assumed safe and live, as well as, in order to guarantee bridge liveness, we assume that at least one Axelar attester is honest (in some cases, this assumption is needed also for bridge safety).

	SPV		Refereed		SNARK
	Sampling	Sync Committee	Linear	Bisection	
<i>Proof Size (Annual)</i>	$365 \cdot 25 \text{ KB}$	$365 \cdot 25 \text{ KB}$	$\frac{365}{x} \cdot (25\text{KB} + x \cdot 32\text{B})$	$\frac{365}{x} \cdot (25\text{KB} + 32\text{B})$	$365 \cdot (209 + 32) \text{ B}$
<i>Non-interactive</i>	✓	✓	✗	✗	✓
<i>Doesn’t rely on Sync Committee</i>	✓	✗	✗	✗	✗
<i>Extra Safety Assumptions</i>	-	-	Exists honest Axelar attester	Exists honest Axelar attester	-
<i>Implementation Complexity</i>	High	Low	Medium	Medium	High
<i>Cost of one prover annually</i>	\$600	\$600	\$600	\$600	>\$500,000

Table 1: Comparison of alternative solutions. All solutions assume that: i) Ethereum is secure (safe and live); ii) to guarantee bridge liveness, there exists an honest Axelar attester. In the refereed case, we assume that the bridge is updated every x days (where 1 sync committee period is equal to approx. 1 day). Also, in the refereed case we compute the best-case scenario, where no disputes occur. In the SNARK case we assume the usage of Groth16 proofs. Our recommendation is highlighted in gray.

Our recommendation

We recommend implementing the second option, that is an SPV client that relies on the sync committee. From the above comparison, it is evident that the first option (SPV client with Ethereum validator sampling) is rather experimental and requires significant research work before proceeding with implementation.

Due to the high traffic that the bridge is expected to observe, it will be updated at least on a daily basis, so the main benefit of the refereed constructions would be lost.

Finally, the exceptional cost of creating SNARK proofs at this point in time makes it prohibitively expensive as an option. Nonetheless, we will design the bridge in a modular way, such that a SNARK solution can be easily incorporated in the future, when proof creation becomes more cost efficient.

3 Construction & Implementation Details

In this section, we discuss the construction and implementation details of the suggested SPV style sync committee based light client. Note that the details of the construction are not final and are subject to change. The client will be implemented as CosmWasm smart contracts. On a high level, the new construction would work as follows:

1. A user initiates the bridge transfer by sending a transaction to a designated smart contract on the Ethereum (source) chain, similar to the existing construction.
2. Once the transaction is finalized, a (possibly different) user submits on the Axelar (destination) chain a light client proof consisting of two components: (i) a consensus verification proof; (ii) an execution verification proof of the corresponding Ethereum transaction.
3. The proofs are verified by the CosmWasm smart contracts on-chain; if valid, the transaction is executed on the Axelar chain.

3.1 Consensus Verification

This module is responsible for keeping track of the latest finalized block and the latest sync committee. To verify the finalized block and the sync committee we use the aggregate sync committee signatures as a proof. Now we will discuss how we can obtain this proof using existing beacon chain APIs and how these proofs can be verified on chain.

- **Update Latest Sync Committee.** This proof validates the sync committee handover, that is updating the committee from one period to the next. It can be generated using Ethereum’s beacon chain API,¹³ specifically the endpoint `/eth/v1/beacon/light_client/updates`. The proof’s verification involves checking the BLS aggregate signature of the sync committee and the Merkle inclusion proof of the next sync committee to the signed header.
- **Update Latest Finalized Block** This proof allows for updating the latest finalized block. It can be generated using the beacon chain API endpoint `/eth/v1/beacon/light_client/finality_update`. Its verification involves checking the BLS aggregate signatures of the latest sync committee on the latest finalized block.

3.2 Execution Verification

This component is responsible for generating and verifying the execution verification proof. We note that the scope of this work includes only verifying if a transaction is included in a block or a certain event was released by a smart contract.

The execution proof mainly constitutes of a Merkle inclusion proof of the relevant transaction or event to the finalized block header received from the consensus verification component. The proof is generated by fetching the complete block, which contains all the transactions and events, and generating the Merkle inclusion proof. Following, the proof verification is a straightforward Merkle inclusion check.

4 Milestones

This project will be implemented in the course of 4 – 6 months and will be executed in multiple phases, described in detail below.

¹³<https://ethereum.github.io/beacon-APIs>

4.1 Phase 1: Architectural Design

During this phase, we will design the architecture of the CosmWasm smart contract implementing the light client, following the Interchain Amplifier's API. The architecture will allow extending the light client to different source chains, beyond Ethereum, and upgrading to different proof mechanisms (e.g., SPV or ZK), making the protocol future-proof.

Deliverables. At the conclusion of this phase, the following items will be delivered:

- A detailed design document, that describes the architecture of the light client to be implemented on Axelar, following Interchain Amplifier standards.

The document will be delivered in LaTeX PDF format and include pseudocode for the relevant components, so that their role is clear from a theoretical point of view.

4.2 Phase 2: Consensus Validation

During this phase, the consensus components of the on-chain light client will be implemented. The implementation will be based on the outcome of Phase 1. During the consensus implementation phase the following will be implemented:

- verification of the Ethereum sync committee *handover signatures*, including aggregate signature verification and quorum logic;
- verification of latest finalized block using sync committee signatures and Ethereum finalization logic.

Deliverables. At the conclusion of this phase, the following items will be delivered:

- The CosmWasm contracts, which implement the consensus components of the on-chain Ethereum light client.
- A test suite, which verifies the correctness of the consensus components.

4.3 Phase 3: Execution Validation

This phase consists of implementing the execution logic of the on-chain Ethereum light client, responsible for checking EVM-related data (such as Ethereum transactions and events). In particular, we will implement the verification of Ethereum execution data, given a validated finalized Ethereum block. This includes the verification of transactions and events and checking the validity of the Merkle proofs for transactions and/or events as needed.

Deliverables. At the conclusion of this phase, the following items will be delivered:

- An extension of the CosmWasm smart contracts (developed during Phase 2), that implements the execution components of the on-chain Ethereum light client.
- A test suite that verifies the correctness of the execution components of the CosmWasm smart contracts.

4.4 Phase 4: Theory, Deployment & Audit

During this phase, the light client's construction will be formalized and its security guarantees will be proven in a formal and rigorous manner. The codebase will be audited and any issues that arise will be fixed.

Deliverables. At the conclusion of this phase, the following items will be delivered:

- Testnet deployment of the Ethereum light client to the Axelar network.
- A LaTeX document that formalizes the light client construction and proves its security guarantees.
- Any fixes to the codebase that may be required as a result of audits.

References

1. S. Agrawal, J. Neu, E. N. Tas, and D. Zindros. Proofs of proof-of-stake with sublinear complexity, 2023.
2. E. Buchman, J. Kwon, and Z. Milosevic. The latest gossip on bft consensus, 2019.
3. V. Buterin et al. A next-generation smart contract and decentralized application platform. *white paper*, 2014.
4. V. Buterin and V. Griffith. Casper the friendly finality gadget. *arXiv preprint arXiv:1710.09437*, 2017.
5. R. Canetti. Universally composable security: A new paradigm for cryptographic protocols. pages 136–145, 2001.
6. Ethereum. Minimal light client, 2023.
7. J. Groth. On the size of pairing-based non-interactive arguments. pages 305–326, 2016.
8. O. Kharif. Key player in ethereum infrastructure infura rejects centralization claim, 2022.
9. P. McCorry, C. Buckland, B. Yee, and D. Song. Sok: Validating bridges as a scaling solution for blockchains. Cryptology ePrint Archive, Paper 2021/1589, 2021. <https://eprint.iacr.org/2021/1589>.
10. A. Network. Axelar network: Connecting applications with blockchain ecosystems, Jan 2021.
11. A. Network. Axelar amplifier: Permissionless connections service run on the axelar network, 2024.
12. E. N. Tas, D. Zindros, L. Yang, and D. Tse. Light clients for lazy blockchains, 2022.
13. Uniswap. Bridge assessment report, 2023.
14. G. Wood. Ethereum: A secure decentralised generalised transaction ledger. *Ethereum Project Yellow Paper*, 151:1–32, 2014.
15. T. Xie, J. Zhang, Z. Cheng, F. Zhang, Y. Zhang, Y. Jia, D. Boneh, and D. Song. zkBridge: Trustless Cross-chain Bridges Made Practical. In *CCS*, pages 3003–3017. ACM, 2022.
16. A. Zamayatin, M. Al-Bassam, D. Zindros, E. Kokoris-Kogias, P. A. Moreno-Sánchez, A. Kiayias, and W. J. Knottenbelt. Sok: Communication across distributed ledgers. In *IACR Cryptology ePrint Archive*, 2019.

5 Appendix

5.1 Theoretical Analysis

In this section we present the model for the theoretical analysis of the bridge framework. We also present conjectures that we plan to prove if the grant is approved.

5.2 Preliminaries

Execution model. A ledger protocol Π is a distributed protocol that offers a *read* and *write* functionality. The ledger protocol is accompanied by a *validity language* \mathbb{V}_Π containing all possible transactions that are “legal” according to the protocol Π . The *read* functionality returns a *ledger* in \mathbb{V}_Π^* , whereas the *write* functionality accepts a ledger and a *transaction* in \mathbb{V}_Π . The *ledger* is a finite sequence of pairs (t, tx) consisting of an integer timestamp t and a transaction tx .

A particular protocol *execution* is one run of an environment $\mathcal{Z}_{\Pi, \mathcal{A}}$ [5] which simulates the execution of Π among multiple honest and adversarial parties. The adversarial parties are controlled by one PPT adversary \mathcal{A} . The adversary and honest parties are modelled as Interactive Turing Machines. Out of the n_Π parties maintaining the ledger, t_Π are corrupt and are controlled by the adversary \mathcal{A} .

Networking. Time evolves in *synchronous* lock-step *rounds* denoted by the integers $1, 2, \dots$. All parties have perfectly synchronized clocks, as they know the current round number. Messages which are broadcast by an honest party at a round r are received by all other honest parties at the beginning of the next round $r + 1$. The adversary can re-order messages, potentially different for each honest party, and inject different messages to the network tapes of different honest parties, but cannot drop messages. We work in the *client gossiping model*, where a new message received by any honest party is rebroadcast

to all other honest parties; therefore, a party cannot identify the origin of a message. We assume there are no bandwidth constraints. The total execution time is polynomial, and each (honest or adversarial) party is bound to polynomial time per round.

Sequence notation. $x \in A$ denotes that either x is an element of set A or x appears in the sequence A . We use $|A|$ for the length of the sequence A . We write $A[i]$ for the i -th element of A (0-based) and $A[-i]$ for the i -th element of the inverse A (1-based); the element $A[-1]$ is the last element of A . We denote by $A[i:j]$ the subarray of A starting from element i (inclusive) and ending in element j exclusive. The notation $A[:j]$ means the sequence from the beginning up to j , whereas $A[i:]$ means the sequence from i till the end.

Ledgers. We denote ${}^{\Pi}L_r^P$ the result of executing the *read* functionality by party P operating in ledger protocol Π at round r . We will write $a \in S$ for an element a and a sequence S if the element a appears somewhere in the sequence S . For a ledger L we will also use the notation $\text{tx} \in L$ to mean that the transaction tx appears in L with some timestamp, that is, there exists some $t \in \mathbb{N}$ such that $(t, \text{tx}) \in L$. In all practical blockchain-based ledger protocols, the timestamp t associated with a transaction will correspond to the timestamp recorded on the block in which the transaction is confirmed. As such, timestamps will all be in the past and non-decreasing.

A ledger protocol is *safe* if the view of different honest parties is consistent.

Definition 1 (Persistence). A ledger protocol Π is persistent during I if for all honest parties P_1, P_2 at rounds $r_1, r_2 \in I$, with $r_1 < r_2$, we have that ${}^{\Pi}L_{r_1}^{P_1} \preceq {}^{\Pi}L_{r_2}^{P_2}$.

A ledger protocol is *live* if honest transactions make it to the ledger.

Definition 2 (Liveness). A ledger protocol Π is live with liveness $u \in \mathbb{N}$ during I if, for all honest parties P_1, P_2 , whenever P_1 attempts to write a valid transaction tx to the ledger at round $r \in I$, then for all $r' \in I$ with $r' \geq r + u$ the transaction is included in ${}^{\Pi}L_{r'}^{P_2}$.

Definition 3 (Timeliness). A ledger protocol Π is timely with timeliness $\nu \in \mathbb{N}$ during I if, for all honest parties P , and for all rounds r , it holds that:

1. ${}^{\Pi}L_r^P[-1]$ has a timestamp in the past, and
2. ${}^{\Pi}L_r^P$ has non-decreasing timestamps.

Additionally, for all rounds $r \in I$, it holds that the timestamps recorded in ${}^{\Pi}L_r^P[| {}^{\Pi}L_{r-1}^P |:]$ are after $r - \nu$.

Ledger security is defined as a ledger protocol that is both safe and live.

Definition 4 (Security). A ledger protocol is secure during I with liveness u if it is:

1. persistent during I ,
2. live with liveness u during I , and
3. timely with timeliness τ during I .

Bridges. A bridge $\Lambda(\Pi_1, \Pi_2)$ is an interoperability protocol between two ledger protocols Π_1 and Π_2 . The execution is defined by a *shared* environment \mathcal{Z} between Π_1 and Π_2 , but each of Π_1 and Π_2 are simulated internally by the environment as before. The purpose of the bridge is to relay events or information that take place on the source side to the destination side. Without loss of generality, we consider Π_1 to be the *source*, and Π_2 to be the *destination*. If the bridge is bidirectional, our statements can be applied in both directions. For our purposes, the purpose is to move value from one side to the other.

A population of n nodes, among which at most t are adversarial, are responsible for operating the bridge. The bridge transmits certain transactions of interest from Π_1 to Π_2 . These are *cross-chain* transactions and are defined by a function ϕ which accompanies the bridge protocol.

Definition 5 (Cross-chain transaction). Let ϕ be an efficiently computable and invertible one-to-one injection $V_{\Pi_1} \rightarrow V_{\Pi_2} \cup \{\perp\}$. A transaction tx on the source side is a cross-chain transaction if $\phi(tx) \neq \perp$.

A bridge is secure (Definition 8) if it guarantees two fundamental properties: safety (Definition 6) and liveness (Definition 7).

Definition 6 (Bridge Safety). A bridge protocol $\Lambda(\Pi_1, \Pi_2)$ is safe with safety $u_s \in \mathbb{N}$ during I if, for all honest parties P_1 of Π_1 and P_2 of Π_2 and for all rounds $r_1, r_2 \in I$ with $r_2 \geq r_1 + u_s$ we have that, whenever $(t, tx) \in {}^{\Pi_2}L_{r_1}^{P_2}$ with $t \in I$ and $\phi^{-1}(tx) \neq \perp$, then $\phi^{-1}(tx) \in {}^{\Pi_1}L_{r_2}^{P_1}$.




Fig. 2: Bridge safety

Definition 7 (Bridge Liveness). A bridge protocol $\Lambda(\Pi_1, \Pi_2)$ is live with liveness $u_\ell \in \mathbb{N}$ during I if, for all honest parties P_1 of Π_1 and P_2 of Π_2 and for all rounds $r_1, r_2 \in I$ with $r_2 \geq r_1 + u_\ell$ we have that, whenever $(t, tx) \in {}^{\Pi_1}L_{r_1}^{P_1}$ with $t \in I$ and $\phi(tx) \neq \perp$, then $\phi(tx) \in {}^{\Pi_2}L_{r_2}^{P_2}$.

Definition 8 (Bridge Security). A bridge Λ is secure with safety $u_s \in \mathbb{N}$ and liveness $u_\ell \in \mathbb{N}$ during I if it is safe with safety u_s during I and live with liveness u_ℓ during I .

Conjecture 1 (Bridge security under temporary dishonest majority). We conjecture that a bridge Λ is safe and live under temporary dishonest majority.




Fig. 3: Bridge liveness

About Common Prefix

Common Prefix is a blockchain research and development company. We believe blockchains will re-organize the world's economy to achieve planetary-scale efficient resource allocation and governance. We do the science and engineering to bring the blockchain field towards mainstream adoption and real-world impact, tackling foundational problems such as scalability, interoperability, and usability. We specialize in all aspects of blockchain science, from the low-level consensus of Layer-1s to the high-level tokenomics of DeFi applications. We have in-house scientist experts in game theory, incentives, auctions, cryptography, zero knowledge, multiparty computation, light clients, wallets, signature schemes, DeFi, and smart contracts from top universities worldwide. We are a team with multichain expertise, with engineers working in ecosystems ranging from Bitcoin and Ethereum, to Cosmos, Cardano, Sui, and Solana. We work with blockchain-first-only industry partners who push the boundaries of what is possible. We help them design, analyze, implement, deploy, and commercialize rigorous protocols from first principles, with provable security and pragmatism in mind.

